

You're your own best teacher: A Self-Supervised Learning Approach For Expressive Representations

Johan Vik Mathisen

June 1, 2024

0.1 Implementation details

In our implementation, we closely follow to the methodology outlined in [TimeVQVAE], particularly in the design and deployment of the encoder, decoder, and codebook.

Time Frequency Modelling

We utilize the Short Time Fourier Transform (STFT) and its inverse (ISTFT) for time-frequency modeling, employing `torch.stft` and `torch.istft` functions, respectively. Consistent with [TimeVQVAE], we set the primary parameter `nfft` to 8 and maintain default settings for others. This configuration yields a frequency axis spanning $[1, 2, 3, 4, 5]$ and a time axis half the length of the original.

Encoder and decoder

The encoder and decoder architectures closely resemble those described in [nadavbh12], with further adaptations from [TimeVQVAE].

As illustrated in Figure ??, the encoder consists of n downsampling convolutional blocks (Conv2d - BatchNorm2d - LeakyReLU), followed by m residual blocks (LeakyReLU - Conv2d - BatchNorm2d - LeakyReLU - Conv2d). Downsampling convolutional layers are implemented with parameters: kernel size=(3,4), stride=(1,2), padding=(1,1), downsampling the temporal axis by a factor of 2 per block. Residual convolutional layers have parameters: kernel size=(3,3), stride=(1,1), padding=(1,1). The decoder mirrors this architecture, featuring m residual and n upsampling layers using transposed convolutional blocks with identical parameters.

The downsampling rate is determined by 2^n , ensuring the resulting latent representation (z) has a width of 32. For a comprehensive understanding of the TimeVQVAE implementation details, we direct readers to Appendix C.3 of [TimeVQVAE].




Figure 1: Overview of the encoder architecture. The decoder architecture is simply obtained by reversing the arrows and switching out the convolutional block for transposed convolutional blocks.

VQ

We adopt a codebook size of 32 and a dimension of 64, utilizing exponential moving average with a decay of 0.9 and a commitment loss with weight $\beta = 1$.

Augmentations

window warp: window ratio: 0.4,min window warp: 0.9, max windowwarp: 2.0
 amplitude resize: Amp Rate: 0.2
 gaussian noise: gaus mean: 0, gaus std: 0.05
 slice and shuffle:

SSL

For both Barlow Twins and VlbCReg, we implement the projector following the guidelines in [lee2024computer]. The projector’s dimension size is set to 4096 for both methods, with λ set to 5×10^3 for Barlow Twins. For VlbCReg, λ and μ are both set to 25, while ν is set to 100.

Prior learning

In the iterative decoding algorithm, the number of iterations T is set to 10, following [chang2022maskgit]. We employ cosine as the mask scheduling function (γ), adopting the implementation from [TimeVQVAE].

Optimizer

We utilize the AdamW optimizer with batch sizes set to 128 for stage 1 and 256 for stage 2, an initial learning rate of 10^{-3} , cosine learning rate scheduler, and a weight decay of 10^{-5} . Both stage 1 and stage 2 training procedures run for 1000 epochs.

Evaluation

For evaluation purposes, K-nearest neighbors (KNN) and Support Vector Machines (SVM) are implemented using scikit-learn, with $K = 5$ for KNN and a linear kernel for SVM.

0.2 Initial Experimentation and Model Development

The overarching objective in creating our model is to learn more expressive latent representations for better time series generation. We want to keep the reconstruction capabilities of the tokenization model at a high level, where the rationality

is that if the tokenization model reconstructs well the latent representations contains all relevant information of the input. We simultaneously want enforce better class separability in the latent representations, as we hypothesize that such additional structure eases/improved learning of the generative model.

During development we encountered several problems:

When we attempted a siamese architecture, with quantization in the augmented branch, and to derive the SSL loss from the discrete representations there were a correlation problem. The codewords were very highly correlated, which resulted from the passing both views through the VQ. $SSL(z_q, z'_q)$

In an attempt to solve this we attempted to derive the SSL loss from the continuous latent representations, but the resulting discrete latent representations performed poorly on the downstream classification task. Separability problem: $SSL(z, z')$

The solution was to remove the VQ in the augmented branch and rather derive the SSL loss from z_q and z' . Solution: $SSL(z_q, z')$

Overfitting problem: Using $SG()$ on augmented branch / Not using augRecons

0.3 Main Experiments

In this thesis, we focus on two main objectives. Firstly, in Stage 1, we aim to determine whether NC-VQVAE can learn more expressive representations compared to VQVAE. Specifically, we investigate whether NC-VQVAE can achieve reconstruction performance on par with VQVAE while simultaneously enhancing downstream classification. In Stage 2, our interest lies in examining the impact of NC-VQVAE on prior learning and time series generation.

To evaluate our model, NC-VQVAE, we compare it against the naive VQVAE using both Barlow Twins and VIBCREg as self-supervised learning (SSL) methods. For each SSL method, we train three distinct models employing different sets of augmentations. Furthermore, each configuration is trained using four different seeds, resulting in a total of 364 models trained at each stage.

Our evaluation process begins with assessing the tokenization models, focusing on their reconstruction capability and performance in downstream classification tasks. Subsequently, we proceed to train a prior model on top of the various tokenization models. We then evaluate the performance of the generative models using metrics such as IS (Inception Score), FID (Fréchet Inception Distance), and CAS (Classification Accuracy Score). Additionally, visual inspections are conducted at both stages to provide further insights into the models' performance.

0.4 Stage 1

One of the most interesting aspects of our model is the choice of augmentations.

0.4.1 Augmentations

In our experiments we consider three sets of augmentations with different characteristics. They are

- Amplitude Resizing + Window Warp
- Slice and Shuffle
- Gaussian noise

Amplitude Resizing + Window Warp scales in both x and y direction. The window warp has similar qualities to phase shift, but not uniformly and keeps endpoints fixed. They were considered as the observed conditional distribution in some datasets, such as ShapesAll ??, had similar overall shape, but peaked with different amplitude and at different locations. Thus the augmented view had similar characteristics as the conditional distribution of the original view ???. In many cases we will refer to this set of augmentations simply as warp.




Figure 2: ShapesAll: Original (left), augmented (right). 15 instances of Amplitude Resizing + Window Warp applied to the original sample.

Slice and Shuffle crops the time series into four sections and permutes them. For datasets with sharp modularity and few peaks, such as ElectricDevices ??, the augmentation provides a view with peaks occurring at timestamps not seen in the training data, which is illustrated in figure ??. This could improve the reconstruction on unseen data, as well as encouraging the model to focus more on the existence of a peak rather than its specific location. For some datasets such as FordA ??, the semantics of the dataset is preserved under this augmentation, despite their continuous nature. In many cases we will refer to this augmentation simply as slice.

Gaussian noise adds a noise $\epsilon \sim N(0, 0.05)$ to each datapoint in the time series. This introduces, in many cases, a substantial high frequency component as seen in figure ??. As the naive VQVAE described in [TimeVQVAE] had trouble with reconstruction of HF components, this augmentation could provide more emphasis on these. The reconstruction of the augmented views can too provide more information regarding HF components for the decoder. Of the three aug-




Figure 3: ElectricDevices: Original (left), augmented (right). 5 instances of Slice and Shuffle applied to the original sample.

mentations, gaussian noise provides the most predictable augmented views from a numerical standpoint, which might result in a SSL loss which is easier to minimize. In many cases we will refer to this augmentation simply as Gaussian or gauss.




Figure 4: StarLightCurves: Original (left), augmented (right). One instance of Gaussian noise applied to the original sample.

0.4.2 Evaluation

The tokenization model, as we are interested in representation learning, is evaluated on two metrics. Firstly, and most importantly its ability to reconstruct the input data once compressed into latent space. In essence the latent representation encodes "everything" (important information is preserved) about the original data if the model is able to reconstruct well. Secondly we evaluate linear classifiers on the latent representations, which provides good results if the model learns discriminative features of the different classes and produces an approximately linear separable space. Finally, as the tokenization model is a part of the generative model, the ultimate evaluation metric is the corresponding evaluation of the generative model.

0.5 Stage 2

0.5.1 Evaluation

- **IS:**
- **FID:**
- **CAS:** We evaluate the CAS for TSTR by using the Supervised FCN on all our models considered and compare against the baseline model to investigate the relative performance.

- Visual inspection:
- Token usage:

0.6 UCR Time Series Classification Archive

The evaluation of our model NC-VQVAE is done on a subset of the UCR Time Series Archive [UCRArchive2018]. The UCR archive is a collection of 128 datasets of univariate time series for time series classification. The different datasets in the archive span a wide range characteristics and include among others sensor, device, image-derived and simulated data. Each dataset has a predefined training and test split.

Our subset of the UCR archive is

Type	Name	Train	Test	Class	Length
Device	ElectricDevices	8926	7711	7	96
Sensor	FordB	3636	810	2	500
Sensor	FordA	3601	1320	2	500
Sensor	Wafer	1000	6164	2	152
Simulated	TwoPatterns	1000	4000	4	128
Sensor	StarLightCurves	1000	8236	3	1024
Motion	UWaveGestureLibraryAll	896	3582	8	945
ECG	ECG5000	500	4500	5	140
Image	ShapesAll	600	600	60	512
Simulated	Mallat	55	2345	8	1024
Image	Symbols	25	995	6	398
Sensor	SonyAIBORobotSurface2	27	953	2	65
Sensor	SonyAIBORobotSurface1	20	601	2	70

Table 1: The subset of the UCR Archive considered for our experiments.

We choose to test on a subset, rather than on the entire UCR Archive, due to computational limitations as well as to more thoroughly investigate the effect of our models and the role of augmentations. The subset is chosen such that they span a wide range of train set sizes, lengths, classes and type, while the class distributions have visually different characteristics which can be seen from table ?? and figure ??.

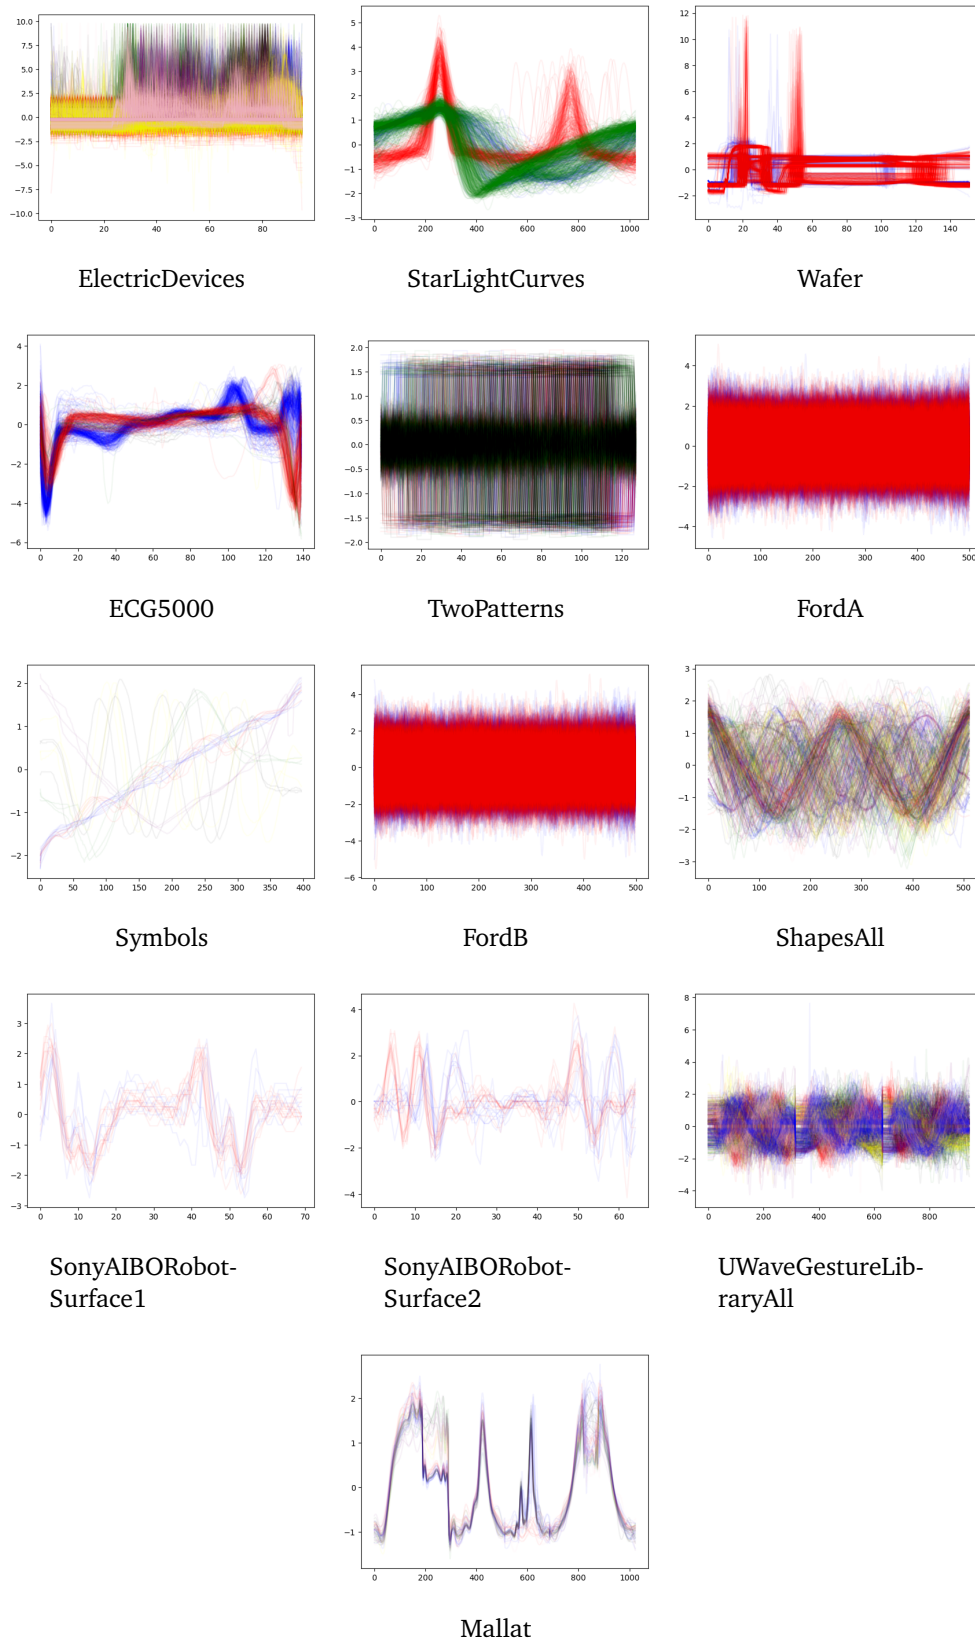


Figure 5: Our selected subset of the UCR Archive. All time series in the training set are plotted and color coded according to label.